

Federated Learning Prototype

by Negin Ebrahimi

Introduction

The objective of this project is to develop a simple FL prototype using PyTorch and the Flower (FL) framework. Both a centralized baseline model and a federated version are implemented and compared in terms of accuracy, convergence, and total training time.

Experiment Setup

- **Dataset:** CIFAR-10 contains 60,000 color images. The dataset was divided into 50,000 training and 10,000 test samples.
- **Model Architecture:** A simple convolutional neural network (CNN) was implemented, consisting of two convolutional layers for feature extraction and two fully connected layers for classification.
- **Federated Learning Configuration:**
 - Clients: 3 clients, each receiving a non-IID (non-identically distributed) subset of CIFAR-10 to simulate realistic heterogeneous data.
 - Aggregation Strategy: FedAvg (Federated Averaging), which averages local client model weights after each round.
 - Training Parameters: Local epochs = 3, Global rounds = 3, Optimizer = Adam, Learning rate = 0.001, Batch size = 64.
 - Hardware: Training was conducted on CPU-only (no GPU acceleration).

Results and Comparison with Centralized Model

Model Type	Accuracy (%)	Loss	Training Time (s)
Centralized	64.28	1.01	~180
Federated (3 clients, 3 rounds)	71.35	0.83	~296

Accuracy over rounds

- Round 1: ~58.6% , Round 2: ~67.8%, Round 3: ~71.3%

This consistent upward trend confirms that the model **converged steadily** with effective weight aggregation.

The **loss decreased** from 1.46 $\rightarrow$ 0.94 $\rightarrow$ 0.83, demonstrating stable optimization and efficient learning.

Convergence Time Analysis

- The communication overhead slightly slowed training, but results remained stable and reproducible.
- Adding more rounds or clients (e.g., 5 clients, 5 rounds) would likely **increase time linearly** but **further improve accuracy**.

Hyperparameter Insights

- **Optimizer:** Adam performed best; SGD led to slower convergence.
- **Learning rate:** 0.001 provided optimal balance; higher rates caused oscillations.
- **Batch size:** 64 balanced accuracy and speed well.
- **Local epochs:** Increasing from 3 $\rightarrow$ 5 slightly improved accuracy but lengthened convergence.

Conclusions

- The **Federated Learning setup** achieved **71.35% accuracy**, outperforming the centralized model's **64.28%**.
- **Convergence time** increased moderately due to multi-client updates.
- Results demonstrate that **FL is effective for decentralized training** while maintaining privacy, making it suitable for distributed scenarios like healthcare or edge devices.

Federated learning converged slightly slower but reached nearly the same final accuracy as the centralized model, confirming the effectiveness of **FedAvg** even under data decentralization.

These results indicate that federated learning can achieve strong global performance without requiring centralized data collection, preserving data privacy while maintaining high model accuracy.